

Statistical Methods for Machine Learning Report

Image Classification with Artificial Neural Networks

Maurina Gabriele
Università degli Studi di Milano

May 15, 2021

Abstract

This document contains a report on the project created by Gabriele Maurina for the exam Statistical Methods for Machine Learning. The project focuses on using artificial neural networks(ANNs) to perform image classification on a dataset from Keggles.com. The dataset contains 90380 images of fruits and vegetables. The project is divided into two experiments. The first experiment trains and evaluates several models with the objective to classify images based on 10 base types, namely "apple", "banana", "cherry", "grape", "peach", "pear", "pepper", "plum", "potato" and "tomato". The second experiment trains and evaluates several models with the objective to classify images based on all 131 labels contained in the dataset. The total number of networks trained and evaluated across both experiments is 300, of which 150 are fully connected deep neural networks (DNNs) and 150 are convolutional neural networks(CNNs). The experimental results show that on average CNNs achieve a lower Zero-One loss on the testing set than DNNs. Furthermore, in the experiment with only 10 labels, CNNs and DNNs achieve similar accuracy and loss, whereas in the experiment with 131 labels, CNNs tend to outperform non-CNNs on accuracy and loss.

Contents

1	Introduction	1
1.1	Deep Neural Networks	2
1.2	Convolutional Neural Networks	2
2	Experiment 1	2
2.1	Setup	2
2.2	Results	3
2.3	Reproducing experiment	3
3	Experiment 2	3
3.1	Setup	3
3.2	Results	3
3.3	Reproducing experiment	3

1 Introduction

Artificial neural networks (ANNs) are machine learning models loosely inspired by biological neural networks found in human and animal brains. ANNs consist of a set of nodes, called artificial neurons, connected to each other by edges. Nodes can send and receive signals, i.e. real numbers, through their edges. Upon receiving a signal from other nodes, a node processes said signal and can signal other nodes. Typically a node output is a function of all its inputs. Nodes can be aggregated into layers, where usually nodes from the first layer forward

signals to nodes in the next layer, which, in turn, forward signals to the next layer and the forwarding is repeated until the last layer is reached.

1.1 Deep Neural Networks

1.2 Convolutional Neural Networks

2 Experiment 1

Experiment 1 trains and evaluates ANNs on the task to classify images of fruits and vegetables of 10 base types. The goal of the experiment is to try different network types in order to find the best configuration, i.e., the configuration with the lowest zero-one loss.

2.1 Setup

This experiment is carried out on a 64 bit machine running Fedora 32, with an Intel[®] Core[™] i7-2670QM CPU @ 2.20GHz and 6GByte of RAM. The programming language used is Python 3.8[4]. Tensorflow2[1], with the Keras[2] interface, is used to create, train and evaluate models. Numpy[3] is used to store and manipulate data easily.

Dataset. The dataset used is available on kaggle.com/moltean/fruits. It contains 90380 images of fruits and vegetables. Experiment 1 classifies images in 10 base types, namely "apple", "banana", "cherry", "grape", "peach", "pear", "pepper", "plum", "potato" and "tomato". Only the 43513 images that represents these fruit and vegetable types are considered. The images are 100x100px photos, already cropped to fit the subject perfectly.

The script `dataset_1.py` is used to manage the dataset. This script performs 5 tasks. 1)It loads images from the dataset. 2)It labels images according to the folder they are in. If the name of the folder starts with one of the 10 labels, the label is attached to the image. Otherwise the image is discarded. 3)It resizes images according to a preferred size. 4)It generates Numpy's n-dimensional arrays containing image data and labels such that they can be easily fed to a machine learning model. The ndarray for image data is of type unsigned 8-bit integer and its shape is `(ni,is,is,3)`, where `ni` is the number of images, `is` is the image side length in pixel, and 3 is the number of channels (r,g,b). The ndarray for labels is of type unsigned 8-bit integer and its shape is `(ni,)`, where `ni` is the number of images. Arrays are created both for training and testing sets. 5)Finally, the script stores said arrays in binary format, so that it will load them faster at the next iteration, without having to go through the whole process again.

The script offers a simple interface, consisting of a single function `dataset(size)`, which takes care of all the aforementioned tasks in the background and returns the dataset, with images of the preferred size, ready to be used by a machine learning model. The `dataset(size)` function is declared as follows:

```

def dataset(size):
    '''Return dataset with images of preferred size.
    If dataset already exists, it is loaded from disk, otherwise it is created
    from the image folder.'''
    files = (
        join(dataset_folder,f'x_train_{size}.npz'),
        join(dataset_folder,f'y_train_{size}.npz'),
        join(dataset_folder,f'x_test_{size}.npz'),
        join(dataset_folder,f'y_test_{size}.npz'))
    for f in files:
        if not isfile(f):
            ds = create_dataset(size)
            save_dataset(*ds,size)
            return ds
    return load_dataset(size)

```

Models. The models used differ in size, shape and type. The script `models.py` handles the creation of models given certain parameters. The parameters supported are depth of the network, i.e., how many layers it has, width of the network, i.e., how many nodes are in each layer, type of network, i.e., DNN or CNN. Furthermore `models.py` can handle different input and output sizes according to the needs.

Execution.

2.2 Results

2.3 Reproducing experiment

3 Experiment 2

Experiment 2, similarly to experiment 1, trains and evaluates ANNs on the task to classify images of fruits and vegetables. Experiment 2, however, considers 131 labels, making the classification task more challenging. The goal of the experiment, as in experiment 1, is to try different network types in order to find the best configuration, i.e., the configuration with the lowest Zero-One loss, and to see if the best configuration for experiment 1, still holds for experiment 2.

3.1 Setup

3.2 Results

3.3 Reproducing experiment

References

- [1] Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu

Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dandelion Mané, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org.

- [2] François Chollet et al. Keras. <https://keras.io>, 2015.
- [3] Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, Robert Kern, Matti Picus, Stephan Hoyer, Marten H. van Kerkwijk, Matthew Brett, Allan Haldane, Jaime Fernández del Río, Mark Wiebe, Pearu Peterson, Pierre Gérard-Marchant, Kevin Sheppard, Tyler Reddy, Warren Weckesser, Hameer Abbasi, Christoph Gohlke, and Travis E. Oliphant. Array programming with NumPy. *Nature*, 585(7825):357–362, September 2020.
- [4] Guido Van Rossum and Fred L. Drake. *Python 3 Reference Manual*. CreateSpace, Scotts Valley, CA, 2009.

Table 1: Experiment 1

IS	T	D	W	ZOL	#E	IS	T	D	W	ZOL	#E	IS	T	D	W	ZOL	#E
10	dense	1	32	0.07	25	20	conv	2	8	0.07	15	40	dense	3	32	0.24	13
10	dense	1	64	0.03	20	20	conv	2	16	0.03	15	40	dense	3	64	0.10	7
10	dense	1	96	0.03	23	20	conv	2	24	0.02	15	40	dense	3	96	0.09	6
10	dense	1	128	0.05	13	20	conv	2	32	0.01	10	40	dense	3	128	0.10	6
10	dense	1	160	0.08	12	20	conv	2	40	0.02	11	40	dense	3	160	0.03	8
10	dense	2	32	0.04	20	20	conv	3	8	0.11	12	40	conv	1	8	0.08	8
10	dense	2	64	0.07	8	20	conv	3	16	0.02	12	40	conv	1	16	0.11	13
10	dense	2	96	0.09	9	20	conv	3	24	0.01	9	40	conv	1	24	0.07	5
10	dense	2	128	0.06	10	20	conv	3	32	0.04	5	40	conv	1	32	0.07	12
10	dense	2	160	0.04	10	20	conv	3	40	0.02	8	40	conv	1	40	0.04	14
10	dense	3	32	0.08	11	30	dense	1	32	0.20	15	40	conv	2	8	0.16	15
10	dense	3	64	0.05	17	30	dense	1	64	0.11	19	40	conv	2	16	0.03	9
10	dense	3	96	0.03	11	30	dense	1	96	0.07	14	40	conv	2	24	0.03	9
10	dense	3	128	0.02	12	30	dense	1	128	0.08	14	40	conv	2	32	0.03	7
10	dense	3	160	0.07	9	30	dense	1	160	0.06	10	40	conv	2	40	0.02	7
10	conv	1	8	0.10	16	30	dense	2	32	0.14	12	40	conv	3	8	0.10	11
10	conv	1	16	0.04	20	30	dense	2	64	0.11	8	40	conv	3	16	0.04	9
10	conv	1	24	0.14	9	30	dense	2	96	0.05	8	40	conv	3	24	0.01	14
10	conv	1	32	0.02	16	30	dense	2	128	0.06	12	40	conv	3	32	0.02	21
10	conv	1	40	0.02	14	30	dense	2	160	0.08	4	40	conv	3	40	0.01	18
10	conv	2	8	0.06	23	30	dense	3	32	0.07	14	50	dense	1	32	0.25	7
10	conv	2	16	0.07	8	30	dense	3	64	0.10	6	50	dense	1	64	0.14	11
10	conv	2	24	0.03	7	30	dense	3	96	0.07	7	50	dense	1	96	0.14	11
10	conv	2	32	0.02	9	30	dense	3	128	0.05	9	50	dense	1	128	0.15	8
10	conv	2	40	0.02	11	30	dense	3	160	0.09	4	50	dense	1	160	0.08	13
10	conv	3	8	0.11	21	30	conv	1	8	0.29	21	50	dense	2	32	0.19	8
10	conv	3	16	0.06	9	30	conv	1	16	0.04	9	50	dense	2	64	0.12	6
10	conv	3	24	0.10	6	30	conv	1	24	0.08	4	50	dense	2	96	0.04	9
10	conv	3	32	0.02	13	30	conv	1	32	0.03	10	50	dense	2	128	0.05	8
10	conv	3	40	0.01	11	30	conv	1	40	0.06	4	50	dense	2	160	0.14	3
20	dense	1	32	0.20	17	30	conv	2	8	0.14	10	50	dense	3	32	0.08	7
20	dense	1	64	0.07	11	30	conv	2	16	0.04	9	50	dense	3	64	0.07	12
20	dense	1	96	0.03	16	30	conv	2	24	0.02	7	50	dense	3	96	0.09	9
20	dense	1	128	0.04	12	30	conv	2	32	0.03	14	50	dense	3	128	0.05	5
20	dense	1	160	0.04	15	30	conv	2	40	0.02	7	50	dense	3	160	0.03	7
20	dense	2	32	0.08	9	30	conv	3	8	0.20	15	50	conv	1	8	0.80	8
20	dense	2	64	0.08	7	30	conv	3	16	0.04	8	50	conv	1	16	0.09	14
20	dense	2	96	0.06	12	30	conv	3	24	0.03	7	50	conv	1	24	0.08	5
20	dense	2	128	0.05	10	30	conv	3	32	0.02	9	50	conv	1	32	0.04	10
20	dense	2	160	0.03	11	30	conv	3	40	0.01	12	50	conv	1	40	0.04	19
20	dense	3	32	0.06	16	40	dense	1	32	0.14	11	50	conv	2	8	0.07	15
20	dense	3	64	0.04	10	40	dense	1	64	0.13	8	50	conv	2	16	0.05	12
20	dense	3	96	0.08	10	40	dense	1	96	0.14	14	50	conv	2	24	0.03	7
20	dense	3	128	0.04	8	40	dense	1	128	0.15	14	50	conv	2	32	0.02	15
20	dense	3	160	0.05	7	40	dense	1	160	0.16	9	50	conv	2	40	0.04	5
20	conv	1	8	0.13	24	40	dense	2	32	0.24	12	50	conv	3	8	0.08	10
20	conv	1	16	0.09	17	40	dense	2	64	0.14	10	50	conv	3	16	0.03	13
20	conv	1	24	0.05	8	40	dense	2	96	0.08	6	50	conv	3	24	0.03	16
20	conv	1	32	0.03	14	40	dense	2	128	0.14	4	50	conv	3	32	0.01	32
20	conv	1	40	0.02	11	40	dense	2	160	0.04	9	50	conv	3	40	0.02	15

IS is image size in pixel;

T is the type of model, can be "dense" or "conv";

D is the depth of the model;

W is the width of the model;

ZOL is the Zero-One loss achieved on the test set;

#E is the number of epochs the model was trained for.

Table 2: Experiment 2

IS	T	D	W	ZOL	#E	IS	T	D	W	ZOL	#E	IS	T	D	W	ZOL	#E
10	dense	1	32	0.07	28	20	conv	2	8	0.12	17	40	dense	3	32	0.38	10
10	dense	1	64	0.08	22	20	conv	2	16	0.09	11	40	dense	3	64	0.18	6
10	dense	1	96	0.08	17	20	conv	2	24	0.14	7	40	dense	3	96	0.10	8
10	dense	1	128	0.08	14	20	conv	2	32	0.05	12	40	dense	3	128	0.12	5
10	dense	1	160	0.07	11	20	conv	2	40	0.09	5	40	dense	3	160	0.16	7
10	dense	2	32	0.09	16	20	conv	3	8	0.24	12	40	conv	1	8	0.44	12
10	dense	2	64	0.09	12	20	conv	3	16	0.08	11	40	conv	1	16	0.07	10
10	dense	2	96	0.07	9	20	conv	3	24	0.10	7	40	conv	1	24	0.99	7
10	dense	2	128	0.09	10	20	conv	3	32	0.12	5	40	conv	1	32	0.08	10
10	dense	2	160	0.07	9	20	conv	3	40	0.06	7	40	conv	1	40	0.14	12
10	dense	3	32	0.10	16	30	dense	1	32	0.95	22	40	conv	2	8	0.13	13
10	dense	3	64	0.11	8	30	dense	1	64	0.37	26	40	conv	2	16	0.19	14
10	dense	3	96	0.08	11	30	dense	1	96	0.30	25	40	conv	2	24	0.12	8
10	dense	3	128	0.12	7	30	dense	1	128	0.13	17	40	conv	2	32	0.08	5
10	dense	3	160	0.13	7	30	dense	1	160	0.13	12	40	conv	2	40	0.04	8
10	conv	1	8	0.30	19	30	dense	2	32	0.34	14	40	conv	3	8	0.13	13
10	conv	1	16	0.11	20	30	dense	2	64	0.16	10	40	conv	3	16	0.22	13
10	conv	1	24	0.11	10	30	dense	2	96	0.08	8	40	conv	3	24	0.07	11
10	conv	1	32	0.06	17	30	dense	2	128	0.11	6	40	conv	3	32	0.07	8
10	conv	1	40	0.07	10	30	dense	2	160	0.11	5	40	conv	3	40	0.05	5
10	conv	2	8	0.19	16	30	dense	3	32	0.55	17	50	dense	1	32	0.99	6
10	conv	2	16	0.12	16	30	dense	3	64	0.14	11	50	dense	1	64	0.99	4
10	conv	2	24	0.09	8	30	dense	3	96	0.12	7	50	dense	1	96	0.53	18
10	conv	2	32	0.10	8	30	dense	3	128	0.08	13	50	dense	1	128	0.99	4
10	conv	2	40	0.10	6	30	dense	3	160	0.13	5	50	dense	1	160	0.53	16
10	conv	3	8	0.21	21	30	conv	1	8	0.17	15	50	dense	2	32	0.99	4
10	conv	3	16	0.11	12	30	conv	1	16	0.68	18	50	dense	2	64	0.91	14
10	conv	3	24	0.09	15	30	conv	1	24	0.67	18	50	dense	2	96	0.89	13
10	conv	3	32	0.14	7	30	conv	1	32	0.27	21	50	dense	2	128	0.31	15
10	conv	3	40	0.05	12	30	conv	1	40	0.14	14	50	dense	2	160	0.16	5
20	dense	1	32	0.52	23	30	conv	2	8	0.21	14	50	dense	3	32	0.90	14
20	dense	1	64	0.31	28	30	conv	2	16	0.10	9	50	dense	3	64	0.56	12
20	dense	1	96	0.12	22	30	conv	2	24	0.04	9	50	dense	3	96	0.16	7
20	dense	1	128	0.09	14	30	conv	2	32	0.07	6	50	dense	3	128	0.14	5
20	dense	1	160	0.09	12	30	conv	2	40	0.07	8	50	dense	3	160	0.15	4
20	dense	2	32	0.31	13	30	conv	3	8	0.40	30	50	conv	1	8	0.99	6
20	dense	2	64	0.08	8	30	conv	3	16	0.13	8	50	conv	1	16	0.39	13
20	dense	2	96	0.09	11	30	conv	3	24	0.08	7	50	conv	1	24	0.41	12
20	dense	2	128	0.10	10	30	conv	3	32	0.09	5	50	conv	1	32	0.21	17
20	dense	2	160	0.14	6	30	conv	3	40	0.06	9	50	conv	1	40	0.15	8
20	dense	3	32	0.13	16	40	dense	1	32	0.94	31	50	conv	2	8	0.39	23
20	dense	3	64	0.11	9	40	dense	1	64	0.81	20	50	conv	2	16	0.27	15
20	dense	3	96	0.12	4	40	dense	1	96	0.30	27	50	conv	2	24	0.22	17
20	dense	3	128	0.09	8	40	dense	1	128	0.41	13	50	conv	2	32	0.11	6
20	dense	3	160	0.11	4	40	dense	1	160	0.24	13	50	conv	2	40	0.06	9
20	conv	1	8	0.80	28	40	dense	2	32	0.99	10	50	conv	3	8	0.49	14
20	conv	1	16	0.37	28	40	dense	2	64	0.27	14	50	conv	3	16	0.10	6
20	conv	1	24	0.10	10	40	dense	2	96	0.11	9	50	conv	3	24	0.07	5
20	conv	1	32	0.11	7	40	dense	2	128	0.11	12	50	conv	3	32	0.07	5
20	conv	1	40	0.07	6	40	dense	2	160	0.18	4	50	conv	3	40	0.05	6

IS is the image size in pixel;

T is the type of model, can be "dense" or "conv";

D is the depth of the model;

W is the width of the model;

ZOL is the Zero-One loss achieved on the test set;

#E is the number of epochs the model was trained for.